

		SRV		NDPipe	
		Fine-tuning	Offline inference	Fine-tuning	Offline inference
Host	Instance	p3.8xlarge		p3.2xlarge	
	CPU	32× vCPUs (2.5GHz)		8× vCPUs (2.3GHz)	
	GPU	2× Tesla V100		1× Tesla V100	
	DRAM	224GB		61GB	
Storage server	Instance	g4dn.4xlarge (GPU disabled)		g4dn.4xlarge	
	CPU	16× vCPUs (2.5GHz)		16× vCPUs (2.5GHz)	
	GPU	-		1× Tesla T4	
	DRAM	64GB		64GB	